

CZ2007

Introduction to Databases

Querying Relational Databases using SQL

Part-2

Arijit Khan

Assistant Professor
School of Computer Science and Engineering
Nanyang Technological University, Singapore

Schedule after Recess Week

SQL

8 Lectures

- Week 8 (Oct 07-Oct 11)
- Week 9 (Oct 14-Oct 18)
- Week 10 (Oct 21-Oct 25)
- Week 11 (Oct 28-Nov 01)

Semi-Structured Data, Quiz-2

2 Lectures

- Week 12 (Nov 02-Nov 08)
- Quiz during Tutorial session
- Quiz syllabus: everything on SQL (Week 8, 9, 10 11)

Summary

- Week 13 (Nov 11-Nov 15)

Recap: Best Practice

- Run your query in the Lab (they usually provide MySQL?)

- (It may not compile, but might still be correct!)

Always check in **Google**



- Consult with the **Book** and course material

- Database Systems: The Complete Book; Hector Garcia-Molina Jeffrey D. Ullman, Jennifer Widom
- (Book available online)

Recap: Roadmap (SQL)

- Introduction to SQL
- Querying single relation



Lecture-1

- **Ordering Tuples**
- **Multi-relation queries**
- **Subqueries**

Lecture-2

Today's lecture: Chapter 6.2, 6.3 of the Book "Database Systems: The Complete Book; Hector Garcia-Molina Jeffrey D. Ullman, Jennifer Widom

- Set operations
- Bag semantics
- Join expressions
- Aggregation

Lectures-3 & 4

Recap: Roadmap (SQL)

- Groupings
- Creation of tables
- Database modifications
- Constraints
- Views

Lecture-5 & 6

- Triggers
- Indexes

Lecture-7 & 8



That would
be all about
Quiz-2!!

Today's Lecture

- Ordering Tuples
- Multi-relation queries
- Subqueries

Study-at-Home slides at the end of every lecture

- They will be in the syllabus of Quiz-2 and Final Exam
- More examples
- Study them at home, will be discussed at the beginning of next lecture
- If any questions, ask me !!



Questions?

The important thing is not to
stop questioning.


Albert Einstein

ORDER BY: Sorting the Results

Product

PName	Price	Category	Manufacturer
Gizmo	19.99	Gadgets	GizmoWorks
Powergizmo	29.99	Gadgets	GizmoWorks
SingleTouch	149.99	Photography	Canon
MultiTouch	203.99	Household	Hitachi

```
SELECT PName, Price, Manufacturer
FROM Product
WHERE Category = 'Gadgets' AND Price < 50
ORDER BY Price, PName
```




PName	Price	Manufacturer
Gizmo	19.99	GizmoWorks
Powergizmo	29.99	GizmoWorks

ORDER BY: Sorting the Results

Product

PName	Price	Category	Manufacturer
Gizmo	19.99	Gadgets	GizmoWorks
Powergizmo	29.99	Gadgets	GizmoWorks
SingleTouch	149.99	Photography	Canon
MultiTouch	203.99	Household	Hitachi

```
SELECT PName, Price, Manufacturer
FROM Product
WHERE Category = 'Gadgets' AND Price < 50
ORDER BY Price, PName
```



- Ordering is ascending, unless you specify the DESC keyword.
- Ties are broken by the second attribute on the ORDER BY list, etc.

PName	Price	Manufacturer
Gizmo	19.99	GizmoWorks
Powergizmo	29.99	GizmoWorks

ORDER BY: Sorting the Results

What about NULL?

NULL is normally treated as less than all non-null values.

- Ordering is ascending, unless you specify the DESC keyword.
- Ties are broken by the second attribute on the ORDER BY list, etc.

Multi-Relation Queries

Sailors


sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

Reserves

sid	bid	day
1	102	9/12
2	102	9/13

Boat reserved
by sailors

```
SELECT S.sname
FROM Sailors S, Reserves R
WHERE S.sid=R.sid AND R.bid=102
```



sname
Fred
Jim

Multi-Relation Queries

Sailors

sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

Reserves


sid	bid	day
1	102	9/12
2	102	9/13

Boat reserved
by sailors

```
SELECT S.sname
FROM Sailors S, Reserves R
WHERE S.sid=R.sid AND R.bid=102
```

```
SELECT S.sname
FROM Sailors AS S, Reserves AS R
WHERE S.sid=R.sid AND R.bid=102
```

Both are OKAY
(Semantically)



sname
Fred
Jim

Multi-Relation Queries

Sailors

sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

Reserves

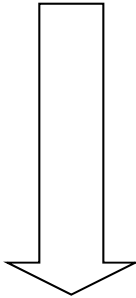
sid	bid	day
1	102	9/12
2	102	9/13

Boat reserved
by sailors

```
SELECT S.sname
FROM   Sailors S, Reserves R
WHERE  S.sid=R.sid AND R.bid=102
```

```
SELECT sname
FROM   Sailors, Reserves
WHERE  Sailors.sid=Reserves.sid AND bid=102
```

Also OKAY
(Semantically)



sname
Fred
Jim


Multi-Relation Queries

Sailors

sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

Reserves

sid	bid	day
1	102	9/12
2	102	9/13



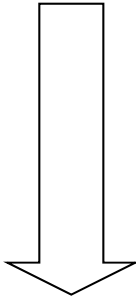
Boat reserved
by sailors

```
SELECT S.sname
FROM Sailors S, Reserves R
WHERE S.sid=R.sid AND R.bid=102
```

```
SELECT Sailors.sname
FROM Sailors, Reserves
WHERE Sailors.sid=Reserves.sid
AND Reserves.bid=102
```



Also OKAY
(Semantically)



sname
Fred
Jim

Multi-Relation Queries (On Same Table)

Sailors

sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

Find pairs of sailor names both of whose ages are below 30



S1.sname	S2.sname
Fred	Nancy

Multi-Relation Queries (On Same Table)

Sailors AS S1


sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

Sailors AS S2

sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

Find pairs of sailor names both of whose ages are below 30

```
SELECT S1.sname, S2.sname
FROM   Sailors S1, Sailor S2
WHERE  S1.age<30 AND S2.age<30
AND    S1.sname < S2.sname
```



S1.sname	S2.sname
Fred	Nancy

Multi-Relation Queries (On Same Table)

Sailors AS S1

sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27


Sailors AS S2

sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

Find pairs of sailor names both of whose ages are below 30

```
SELECT S1.sname, S2.sname
FROM   Sailors S1, Sailor S2
WHERE  S1.age<30 AND S2.age<30
AND    S1.sname < S2.sname
```

String comparison to
avoid duplicates



S1.sname	S2.sname
Fred	Nancy

Questions?

Today's Lecture

- Ordering Tuples



- Multi-relation queries



- Subqueries

Subqueries

SELECT Clause

FROM Clause

SQL

WHERE Clause

SQL

Subqueries

- Also called nested queries
- We can do nested queries because SQL is compositional:
 - Everything (inputs / outputs) is represented as multisets- the output of one query can thus be used as the input to another (nesting)!
- This is extremely powerful!

Types of Subqueries

Scalar Subquery

- returns a **single value** which is then used in a comparison.
- if query expects a single value from a subquery, and it returns multiple values or no values, a **run-time error** occurs.

Row Subquery

- returns a **single row** which may have multiple columns

Table Subquery

- returns **one or more columns and multiple rows**.

Scalar Subquery

Query

From **Sells**(bar, beer, price), find the bars that serve **Heineken** for the same price **WOOBAR** charges for **Bud**.

```
SELECT    bar
FROM      Sells
WHERE     beer = 'Heineken'
AND price = [price of Bud @ WOOBAR];
```

AND price = [price of Bud @ WOOBAR];


Subqueries

- Find the price **WOOBAR** charges for **Bud**.
- Find the bars that serve **Heineken** at that price.

Scalar Subquery

Sells

<u>Bar</u>	<u>Beer</u>	Price
Southbridge	Heineken	7.90
Southbridge	Bud	6.60
WOOBAR	Bud	7.90
WOOBAR	Heineken	8.10
Emerald Hill	Heineken	8.00



Price
7.90

```


SELECT price
FROM Sells
WHERE bar = 'WOOBAR'
AND beer = 'Bud';
  
```

AND beer = 'Bud';

Scalar Subquery

Sells

<u>Bar</u>	<u>Beer</u>	Price
Southbridge	Heineken	7.90
Southbridge	Bud	6.60
WOOBAR	Bud	7.90
WOOBAR	Heineken	8.10
Emerald Hill	Heineken	8.00



Bar
Southbridge

SELECT
FROM
WHERE

price
Sells
bar = 'V'
AND be

SELECT
FROM
WHERE

bar
Sells
beer = 'Heineken'
AND price = 7.90;

Scalar Subquery

```
SELECT    bar
FROM      Sells
WHERE     beer = 'Heineken' AND
          price = ( SELECT price
                    FROM   Sells
                    WHERE  bar = 'WOOBAR'
                    AND    beer = 'Bud');
```

Without using Scalar Subquery?

```
SELECT      S1.bar
FROM        Sells S1, Sells S2
WHERE       S1.beer = 'Heineken'
AND         S2.bar = 'WOOBAR'
AND         S2.beer = 'Bud'
AND         S1.price = S2.price;
```



Use two copies of the table

Row Subquery

Row Subquery

- returns a **single row** which may have multiple columns

Operators in Row Subquery

IN

<tuple> IN <relation> is true if and only if the tuple is a member of the relation.

Row Subquery

Row Subquery

- returns a **single row** which may have multiple columns

Operators in Row Subquery

ALL

$x \neq \text{ALL}(\langle \text{relation} \rangle)$ is true if and only if for every tuple t in the relation, x is **not equal** to t .

Row Subquery

Row Subquery

- returns a **single row** which may have multiple columns

Operators in Row Subquery

ANY/SOME

- $x = \text{SOME}(\text{<relation>})$ is a Boolean condition. Meaning that **x equals** at least one tuple in the relation.
- “Equal to at least one”
- Early version of SQL allowed **ANY**

“IN” - Row Subquery

IN

<tuple> IN <relation> is true if and only if the tuple is a member of the relation.

Query

From **Beers(name, manf)** and **Likes(drinker, beer)**, find the name and manufacturer of each beer that Fred likes.

```
SELECT      *  
FROM        Beers  
WHERE       name IN [What Fred likes]
```

“IN” - Row Subquery

IN

<tuple> IN <relation> is true if and only if the tuple is a member of the relation.

Query

From **Beers(name, manf)** and **Likes(drinker, beer)**, find the name and manufacturer of each beer that Fred likes.

```
SELECT      *
FROM        Beers
WHERE       name IN ( SELECT      beer
                        FROM        Likes
                        WHERE       drinker = 'Fred');
```


Without using Row Subquery?

Query

From Beers(name, manf) and Likes(drinker, beer), find the name and manufacturer of each beer that Fred likes.

```
SELECT      name, manf
FROM        Beers, Likes
WHERE       name = beer
AND         drinker = 'Fred';
```

```
AND         drinker = 'Fred';
```

“ALL” - Row Subquery


ALL

$x \neq \text{ALL}(\text{<relation>})$ is true if and only if for every tuple t in the relation, x is not equal to t .

Query

From $\text{Sells}(\text{bar}, \text{beer}, \text{price})$, find the beer(s) sold for the highest price.

```
SELECT    beer
FROM      Sells
WHERE     price = [highest price];
```



How to find
highest price?

“ALL” - Row Subquery

ALL

$x \neq \text{ALL}(\text{relation})$ is true if and only if for every tuple t in the relation, x is not equal to t .

Query

From $\text{Sells}(\text{bar}, \text{beer}, \text{price})$, find the beer(s) sold for the highest price.

```
SELECT      beer
FROM        Sells
WHERE       price >= ALL ( SELECT      price
                           FROM        Sells );
```

“SOME” - Row Subquery

ANY/SOME

- $x = \text{SOME}(<\text{relation}>)$ is a boolean cond. Meaning that x equals at least one tuple in the relation.
- “Equal to at least one”
- Early version of SQL allowed **ANY**

Query

From Agents(agent_code, agent_name), Customer(agent_code, cust_country), report all agents who belong to the country 'UK'.

“SOME” - Row Subquery

ANY/SOME

- $x = \text{SOME}(<\text{relation}>)$ is a boolean cond. Meaning that x equals at least one tuple in the relation.
- “Equal to at least one”
- Early version of SQL allowed **ANY**

Query

From Agents(agent_code, agent_name), Customer(agent_code, cust_country), report all agents who belong to the country 'UK'.

```
SELECT agent_code, agent_name
FROM Agents
WHERE agent_code = SOME (
    SELECT agent_code FROM Customer
    WHERE cust_country = 'UK');
```

More Operators for Subquery

- Any of the comparison operators (<, <=, =, etc.) can be used.
- The keyword **NOT** can proceed any of the operators (**s** NOT IN **R**)

Table Subquery

Table Subquery

- returns **one or more columns and multiple rows**.

Operators in Table Subquery: Exists/ No Exists

Table Subquery

Table Subquery

- returns **one or more columns and multiple rows**.

Operators in Table Subquery: Exists/ No Exists

Product(name, price, category, maker)

```
SELECT p1.name
FROM   Product p1
WHERE  p1.maker = 'Gizmo-Works'
AND EXISTS(
    SELECT *
    FROM   Product p2
    WHERE  p2.maker <> 'Gizmo-Works'
    AND    p1.name = p2.name)
```

Find products made
by “Gizmo-Works”
having the same
names as products
made by other
makers

Questions?

Summary

- Ordering Tuples



- Multi-relation queries



- Subqueries



Study-at-Home

Correlated and uncorrelated subqueries (Slides 42-44)

Will be in the syllabus of
Quiz-2 and Final Exam

Uncorrelated Subqueries

```
SELECT      *  
FROM        Beers  
WHERE       name IN ( SELECT      beer  
                        FROM        Likes  
                        WHERE       drinker = 'Fred');
```

- Subquery is not related to the outer query

Correlated Subqueries

```
SELECT p1.name
FROM   Product p1
WHERE  p1.maker = 'Gizmo-Works'
AND EXISTS(
    SELECT *
    FROM   Product p2
    WHERE  p2.maker <> 'Gizmo-Works'
    AND    p1.name = p2.name)
```

- A subquery is **correlated** with the outer query if it contains a reference to an attribute in the outer query.
- A subquery is **correlated** with the outside query if it must be **re-computed for every tuple** produced by the outside query.

Subquery – Rules to Remember

- The **ORDER BY** clause may not be used in a subquery.
- The number of attributes in the **SELECT** clause in the subquery must match the number of attributes with the comparison operator.
- Column names in a subquery refer to the table name in the **FROM** clause of the subquery by default.
- When the result of a subquery is used as an operand, it must be the right operand.

Questions ??



Thank You !