

CZ2007

Introduction to Databases

Querying Relational Databases using SQL

Part-3

Arijit Khan

Assistant Professor
School of Computer Science and Engineering
Nanyang Technological University, Singapore

Schedule after Recess Week

SQL

8 Lectures

- Week 8 (Oct 07-Oct 11)
- Week 9 (Oct 14-Oct 18)
- Week 10 (Oct 21-Oct 25)
- Week 11 (Oct 28-Nov 01)

Semi-Structured Data, Quiz-2

2 Lectures

- Week 12 (Nov 02-Nov 08)
- Quiz during Tutorial session
- Quiz syllabus: everything on SQL (Week 8, 9, 10 11)

Summary

- Week 13 (Nov 11-Nov 15)

Recap: Roadmap (SQL)

- Introduction to SQL
- Querying single relation



Lecture-1

- Ordering Tuples
- Multi-relation queries
- Subqueries



Lecture-2

- **Set operations**
- **Bag semantics**
- **Join expressions**
- **Aggregation**

Lectures-3 & 4

Today's lecture: Chapter 6.3, 6.4 of the Book "Database Systems: The Complete Book; Hector Garcia-Molina Jeffrey D. Ullman, Jennifer Widom

Recap: Roadmap (SQL)

- Groupings
- Creation of tables
- Database modifications
- Constraints
- Views

Lecture-5 & 6

- Triggers
- Indexes

Lecture-7 & 8



That would
be all about
Quiz-2!!

Today's Lecture

- Set operations
- Bag semantics
- Join expressions
- Aggregation

Study-at-Home slides at the end of every lecture

- They will be in the syllabus of Quiz-2 and Final Exam
- More examples and cases
- Study them at home
- If any questions, ask me !!



Questions?

The important thing is not to
stop questioning.

Albert Einstein

Set Operations in SQL: UNION, INTERSECT, EXCEPT

- They are generally used to combine the results of two separate SQL queries.
- UNION, INTERSECT, EXCEPT

Syntax

- (subquery) UNION (subquery)
- (subquery) INTERSECT (subquery)
- (subquery) EXCEPT (subquery)


INTERSECT

```
SELECT R.A  
FROM R, S  
WHERE R.A=S.A
```

INTERSECT

```
SELECT R.A  
FROM R, T  
WHERE R.A=T.A
```

$$\{r.A \mid r.A = s.A\} \cap \{r.A \mid r.A = t.A\}$$



INTERSECT

Company(name, hq_city)
Product(pname, maker, factory_loc)

```
SELECT hq_city  
FROM Company, Product  
WHERE maker = name AND factory_loc = 'US'
```

INTERSECT

```
SELECT hq_city  
FROM Company, Product  
WHERE maker = name AND factory_loc = 'China'
```

*“Find
Headquarters of
companies which
make gizmos in
US **AND** China”*



This is incorrect. **What is wrong?**

INTERSECT

Company(name, hq_city)
AS C

Product(pname, maker, factory_loc)
AS P

```
SELECT hq_city  
FROM Company, Product  
WHERE maker = name  
AND factory_loc='US'
```

INTERSECT

```
SELECT hq_city  
FROM Company, Product  
WHERE maker = name  
AND factory_loc='China'
```

X Co. has a factory in the US (but not China)

Y Inc. has a factory in China (but not US)

INTERSECT

Company(name, hq_city)
 AS C

Product(pname, maker, factory_loc)
 AS P

```
SELECT hq_city
FROM Company, Product
WHERE maker = name
AND factory_loc='US'
```

INTERSECT

```
SELECT hq_city
FROM Company, Product
WHERE maker = name
AND factory_loc='China'
```

X Co. has a factory in the US (but not China)

Y Inc. has a factory in China (but not US)

But Seattle is returned by the query!

C.name	C.hq_city	P.pname	P.maker	P.factory_loc
X Co.	Seattle	X	X Co.	U.S.
Y Inc.	Seattle	Y	Y Inc.	China

INTERSECT

Company(name, hq_city)
AS C

Product(pname, maker, factory_loc)
AS P

```
SELECT hq_city
FROM Company, Product
WHERE maker = name
AND factory_loc='US'
```

INTERSECT

```
SELECT hq_city
FROM Company, Product
WHERE maker = name
AND factory_loc='China'
```

X Co. has a factory in the US (but not China)

Y Inc. has a factory in China (but not US)

But Seattle is returned by the query!

We did the INTERSECT on the wrong attributes!

C.name	C.hq_city	P.pname	P.maker	P.factory_loc
X Co.	Seattle	X	X Co.	U.S.
Y Inc.	Seattle	Y	Y Inc.	China

INTERSECT

```
Company(name, hq_city)  
Product(pname, maker, factory_loc)
```

```
SELECT hq_city, name  
FROM Company, Product  
WHERE maker = name AND factory_loc = 'US'
```

INTERSECT

```
SELECT hq_city, name  
FROM Company, Product  
WHERE maker = name AND factory_loc = 'China'
```

*“Find
Headquarters of
companies which
make gizmos in
US **AND** China”*

This is okay.

[But, the output also contains “name” in addition to “hq_city”]

Solution – SELECT INTO

Company(name, hq_city)
Product(pname, maker, factory_loc)

```
SELECT hq_city, name  
INTO HQ_Name  
FROM Company, Product  
WHERE maker = name AND factory_loc = 'US'
```

INTERSECT

```
SELECT hq_city, name  
FROM Company, Product  
WHERE maker = name AND factory_loc = 'China';
```

```
SELECT DISTINCT hq_city  
FROM HQ_Name;
```

*“Find
Headquarters of
companies which
make gizmos in
US **AND** China”*



This is the solution – but it requires **two** SQL queries.

Solution – SELECT INTO

Company(name, hq_city)
Product(pname, maker, factory_loc)

```
SELECT hq_city, name  
INTO HQ_Name  
FROM Company, Product  
WHERE maker = name AND factory_loc = 'US'
```

INTERSECT

```
SELECT hq_city, name  
FROM Company, Product  
WHERE maker = name AND factory_loc = 'China';
```

```
SELECT DISTINCT hq_city  
FROM HQ_Name;
```

*“Find
Headquarters of
companies which
make gizmos in
US **AND** China”*

- SELECT INTO
creates a new
table.

- A physical table
is created

This is the solution – but it requires **two** SQL queries.

Alternative Solution using Subquery

Company(name, hq_city)
Product(pname, maker, factory_loc)

```
SELECT DISTINCT hq_city
FROM Company, Product
WHERE maker = name
      AND name IN (
        SELECT maker
        FROM Product
        WHERE factory_loc = 'US')
      AND name IN (
        SELECT maker
        FROM Product
        WHERE factory_loc = 'China')
```

*“Headquarters of
companies which
make gizmos in US
AND China”*




Union

```
SELECT R.A  
FROM R, S  
WHERE R.A=S.A
```

UNION

```
SELECT R.A  
FROM R, T  
WHERE R.A=T.A
```

$$\{r.A \mid r.A = s.A\} \cup \{r.A \mid r.A = t.A\}$$




Except

```
SELECT R.A  
FROM R, S  
WHERE R.A=S.A
```

EXCEPT

```
SELECT R.A  
FROM R, T  
WHERE R.A=T.A
```

$$\{r.A \mid r.A = s.A\} \setminus \{r.A \mid r.A = t.A\}$$



More Example: Union

From relations Likes(drinker, beer), Sells(bar, beer, price), and Frequents(drinker, bar), find the drinkers and beers such that:

- The drinker likes the beer,
- or
- The drinker frequents at least one bar that sells the beer.

```
(SELECT *  
FROM Likes)  
  
UNION  
  
( SELECT drinker, beer  
FROM Sells, Frequent  
WHERE Frequent.bar = Sells.bar);
```

Questions?

Bag Semantics vs. Set Semantics

- Set semantics → No duplicates, each item appears only once
- Bag semantics → Duplicates allowed, i.e., a multiset

- Default for **SELECT-FROM-WHERE** is **bag**
- Default for **UNION, INTERSECT, and EXCEPT** is **set**

How to change the default?


- Force set semantics with **DISTINCT** after **SELECT**
- Force bag semantics with **ALL** after **UNION, etc.**

DISTINCT: Change Bag Semantics to Set Semantics


```
SELECT DISTINCT Category  
FROM Product
```

Versus

```
SELECT Category  
FROM Product
```



Category
Gadgets
Photography
Household



Category
Gadgets
Gadgets
Photography
Household

ALL: Change Set Semantics to BAG Semantics

Likes

<u>Drinker</u>	<u>Beer</u>
Sally	Heineken
Sean	Bud
Melissa	Tiger

Buys


<u>Drinker</u>	<u>Beer</u>
Sally	Heineken
Sally	Bud
Melissa	Heineken
Melissa	Tiger

```
(SELECT *
FROM Likes)
UNION ALL
(SELECT *
FROM Buys);
```

<u>Drinker</u>	<u>Beer</u>
Sally	Heineken
Sally	Bud
Melissa	Heineken
Melissa	Tiger
Sally	Heineken
Sean	Bud
Melissa	Tiger

Join (\bowtie)

- Joins multiple tables
- Already did some examples while answering queries from multiple tables



Join (\bowtie)

Product(PName, Price, Category, Manufacturer)

Company(CName, StockPrice, Country)

Example: Find all products under \$200 manufactured in Japan; return their names and prices.

```
SELECT PName, Price
FROM Product, Company
WHERE Manufacturer = CName
      AND Country='Japan'
      AND Price <= 200
```

```
SELECT PName, Price
FROM Product
JOIN Company ON Manufacturer = CName
              AND Country='Japan'
WHERE Price <= 200
```

Several equivalent ways to write a basic join in SQL.

A few more ways later on...


Join (⋈) - Example

Product

PName	Price	Category	Manuf
Gizmo	\$19	Gadgets	GWorks
Powergizmo	\$29	Gadgets	GWorks
SingleTouch	\$149	Photography	Canon
MultiTouch	\$203	Household	Hitachi

Company

Cname	Stock	Country
GWorks	25	USA
Canon	65	Japan
Hitachi	15	Japan



```

SELECT PName, Price
FROM Product, Company
WHERE Manufacturer = CName
      AND Country='Japan'
      AND Price <= 200
  
```

PName	Price
SingleTouch	\$149.99

Meaning (Semantics) of Join

```
SELECT x1.a1, x1.a2, ..., xn.ak  
FROM R1 AS x1, R2 AS x2, ..., Rn AS xn  
WHERE Conditions(x1, ..., xn)
```

```
Answer = {}  
for x1 in R1 do  
  for x2 in R2 do  
    .....  
    for xn in Rn do  
      if Conditions(x1, ..., xn)  
        then Answer = Answer  $\cup$  {(x1.a1, x1.a2, ..., xn.ak)}  
return Answer
```

Note: this is a *multiset* union
(bag semantics)

Meaning (Semantics) of Join

– An Example

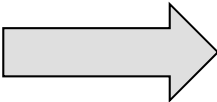
R

A
1
3

```

SELECT R.A
FROM R, S
WHERE R.A = S.B
  
```

Output



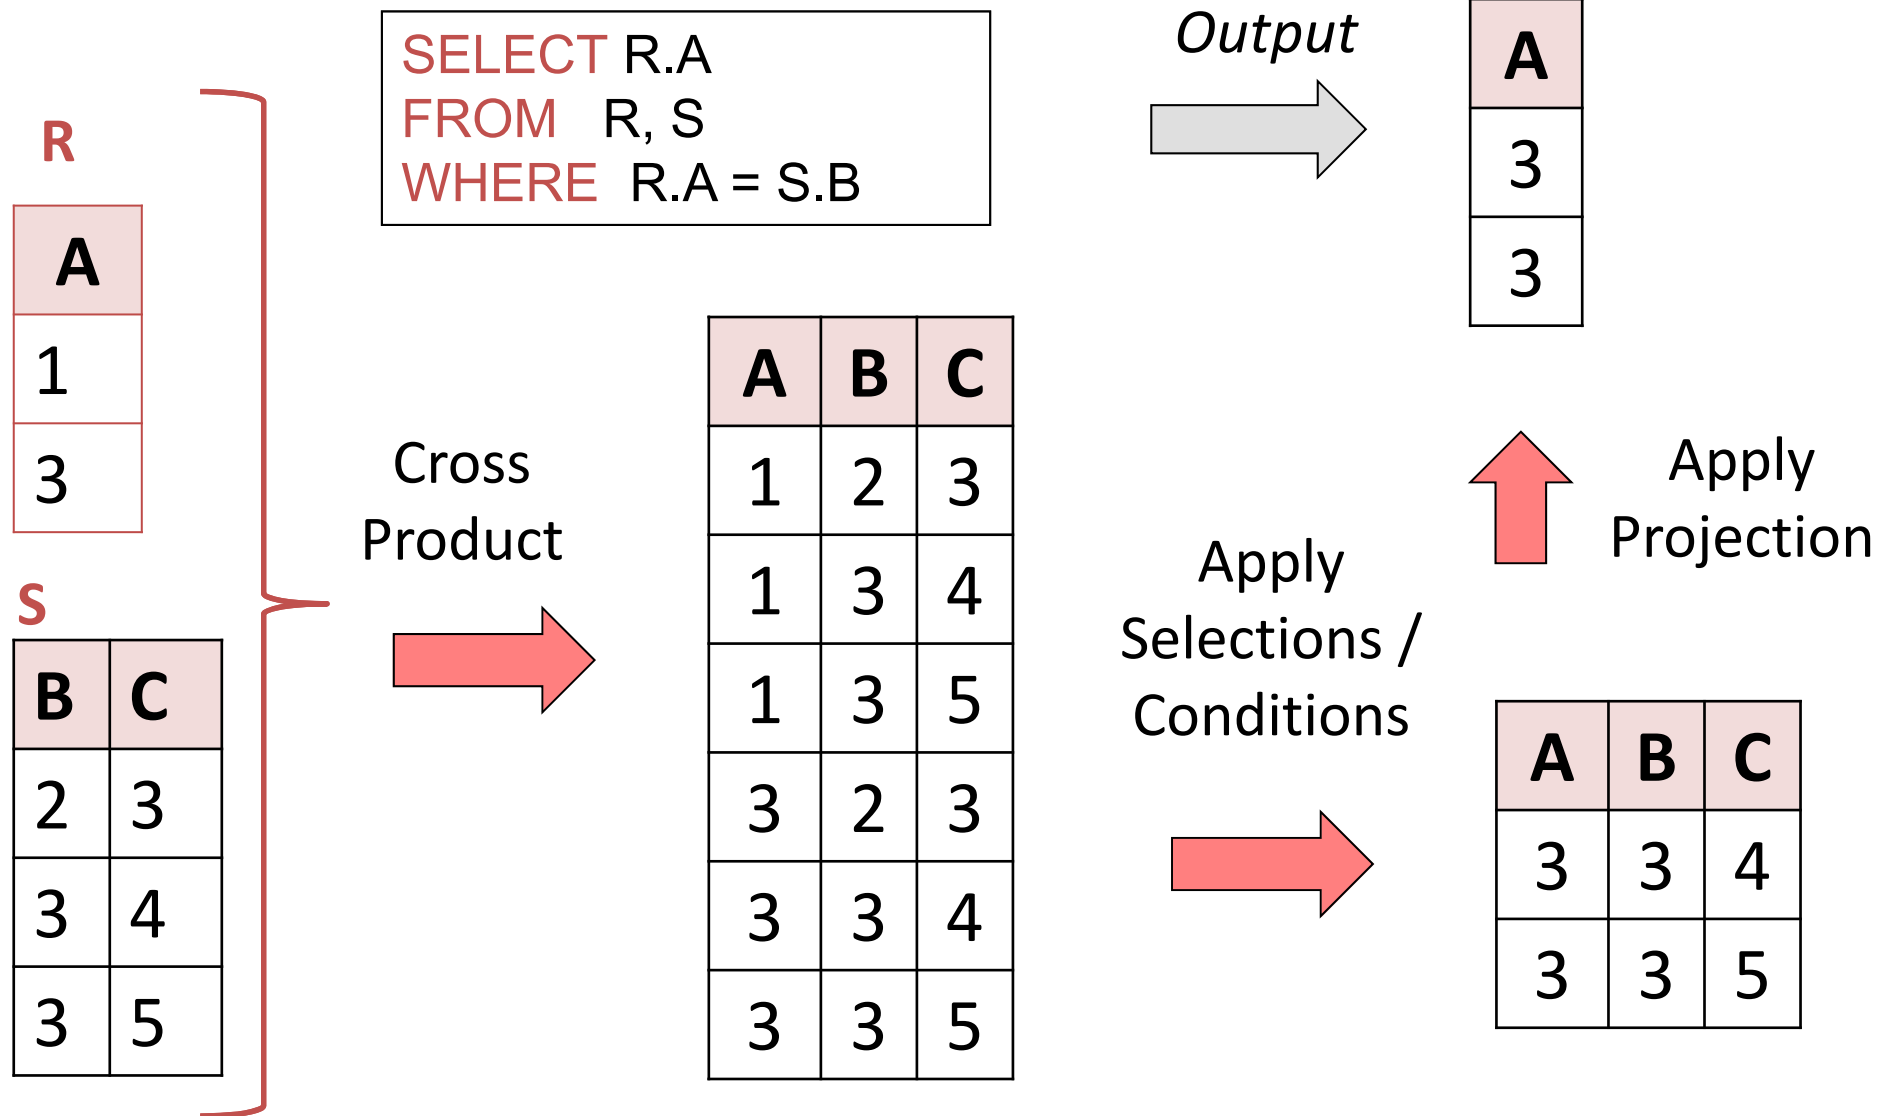
A
3
3

S

B	C
2	3
3	4
3	5

Meaning (Semantics) of Join

– An Example



Meaning (Semantics) of Join

```
SELECT R.A  
FROM R, S  
WHERE R.A = S.B
```

1. Take **cross product**:

$$X = R \times S$$

Recall: Cross product ($A \times B$) is the set of all unique tuples in A, B

Ex: $\{a, b, c\} \times \{1, 2\}$
 $= \{(a, 1), (a, 2), (b, 1), (b, 2), (c, 1), (c, 2)\}$

2. Apply **selections / conditions**:

$$Y = \{(r, s) \in X \mid r.A == s.B\}$$

= Filtering!

3. Apply **projections** to get final output:

$$Z = (y.A,) \text{ for } y \in Y$$

= Returning only *some* attributes
(Bag semantics)

Remembering this order is critical

How Join is Actually Executed in a Database System?

- The preceding slides show what a join means (i.e., semantics)
- Not actually how the DBMS executes it under the covers

We shall not study it in this course
– will be discussed in CZ 4031

Join – More Examples

Product(PName, Price, Category, Manufacturer)

Company(CName, StockPrice, Country)

Find all countries that manufacture some product in the 'Gadgets' category.

```
SELECT Country
FROM Product, Company
WHERE Manufacturer=CName AND Category='Gadgets'
```



Join – More Examples

Product(PName, Price, Category, Manufacturer)

Company(CName, StockPrice, Country)

Find all countries that manufacture some product in the 'Gadgets' category.

```
SELECT Country
FROM Product, Company
WHERE Manufacturer=CName AND Category='Gadgets'
```



```
SELECT DISTINCT Country
FROM Product, Company
WHERE Manufacturer=CName AND Category='Gadgets'
```



Join – A Difficult Example

```
SELECT DISTINCT R.A  
FROM R, S, T  
WHERE R.A=S.A OR R.A=T.A
```

What if $S = \phi$?

```
SELECT DISTINCT R.A  
FROM R, T  
WHERE R.A=T.A
```



Go back to the Join semantics! – the correct answer is ϕ



Join – A Difficult Example

```
SELECT DISTINCT R.A  
FROM R, S, T  
WHERE R.A=S.A OR R.A=T.A
```

- Recall the semantics!
 1. Take cross-product
 2. Apply selections / conditions
 3. Apply projection
- If $S = \{\}$, then the cross product of $R, S, T = \{\}$,
and the query result = $\{\}$!



Join – A Difficult Example

```
SELECT DISTINCT R.A  
FROM R, S, T  
WHERE R.A=S.A OR R.A=T.A
```

- Recall the semantics!
 1. Take cross-product
 2. Apply selections / conditions
 3. Apply projection

```
output = {}  
  
for r in R:  
    for s in S:  
        for t in T:  
            if r['A'] == s['A'] or r['A'] == t['A']:  
                output.add(r['A'])  
return list(output)
```

- If $S = \{\}$, then the cross product of $R, S, T = \{\}$,
and the query result = $\{\}$!



Questions?

- Set operations
- Bag semantics
- Join expressions
- **Aggregation**



SQL Aggregates

```
SELECT AVG(price)
FROM Product
WHERE maker = "Toyota"
```

```
SELECT COUNT(*)
FROM Product
WHERE year > 1995
```

- SQL supports several **aggregation** operations:
 - SUM, COUNT, MIN, MAX, AVG

*Except COUNT, all aggregations
apply to a single attribute*

SQL Aggregates

- COUNT applies to duplicates, unless otherwise stated

```
SELECT COUNT(category)
FROM Product
WHERE year > 1995
```

We probably want:

```
SELECT COUNT(DISTINCT category)
FROM Product
WHERE year > 1995
```

SQL Aggregates

More Rules

- **COUNT**, **MAX**, and **MIN** apply to all types of fields
- **SUM** and **AVG** apply to only numeric fields.
- Except for **COUNT(*)** all functions ignore nulls.
- **COUNT(*)** returns the number of rows in the table.
- Use **DISTINCT** to eliminate duplicates.

More Examples on COUNT

Table

Beer(beer, manufacturer)

```
SELECT      COUNT(manufacturer)
FROM        Beer
```

NULL manufacturers
will be ignored

- Duplicate
manufacturers will be
counted

```
SELECT      COUNT(*)
FROM        Beer
```

NULL manufacturers
will be counted

More Examples on COUNT

Table

Beer(beer, manufacturer)

```
SELECT      COUNT(DISTINCT manufacturer)
FROM        Beer;
```

- Number of distinct manufacturers
- Nulls are ignored

```
SELECT      DISTINCT COUNT(manufacturer)
FROM        Beer;
```

- Number of not-null manufacturers
- Nulls are ignored

Questions?

Summary

- Set operations



- Bag semantics



- Join expressions



- Aggregation



Study-at-Home

Types of SQL Joins (Slides 44-50)

Will be in the syllabus of
Quiz-2 and Final Exam

Types of SQL Join

- Theta Join
- Inner Join
- Natural Join
- Left Outer Join
- Right Outer Join
- Full Outer Join
- (Equi-Join, Product Join, Semi-Join, etc.)

Theta Join

Syntax

- **R JOIN S ON <condition>**
- A theta-join using **<condition>** for selection.

Example

Product(PName, Price, Category, Manufacturer)
Company(CName, StockPrice, Country)

Example: Find all products manufactured in Japan, and stock price more than \$300; return their names and prices.

```
SELECT PName, Price
FROM   Product
JOIN   Company ON Manufacturer = Cname AND Country='Japan'
                AND StockPrice >= 300
```

Theta Join

Syntax

- **R JOIN S ON <condition>**
- A theta-join using **<condition>** for selection.

Example

Product(PName, Price, Category, Manufacturer)
Company(CName, StockPrice, Country)

Example: Find all products manufactured in Japan, and stock price more than \$300; return their names and prices.

```
SELECT PName, Price
FROM Product
JOIN Company ON Manufacturer = Cname AND Country = 'Japan'
AND StockPrice >= 300
```

Any Boolean
condition

Inner Join

Syntax

- `R INNER JOIN S USING (<attribute list>)`
- `R INNER JOIN S ON R.column_name = S.column_name`

Example

TableA

Column1	Column2
1	2

TableB

Column1	Column3
1	3

The INNER JOIN of **TableA** and **TableB** on Column1 will return:

TableA.Column1	TableA.Column2	TableB.Column1	TableB.Column3
1	2	1	3

```
SELECT * FROM TableA INNER JOIN TableB USING (Column1)
```

```
SELECT * FROM TableA INNER JOIN TableB ON TableA.Column1 = TableB.Column1
```


Natural Join

Syntax

- **R NATURAL JOIN S**

Example

TableA

Column1	Column2
1	2

TableB

Column1	Column3
1	3

The NATURAL JOIN of **TableA** and **TableB** will return:

Column1	Column2	Column3
1	2	3

SELECT * FROM TableA NATURAL JOIN TableB

- The repeated columns are avoided.
- One can not specify the joining columns in a natural join.

Outer Join

Syntax

- **R OUTER JOIN S** is the core of an outerjoin expression.

Different Variants

- Optional **NATURAL** in front of **OUTER**.
- Optional **ON <condition>** after **JOIN**.
- Optional **LEFT, RIGHT, or FULL** before **OUTER**.
 - **LEFT** = pad dangling tuples of **R** only.
 - **RIGHT** = pad dangling tuples of **S** only.
 - **FULL** = pad both; this choice is the default.

Example

- `Loan(loanNo,branch,amount), Borrower(cName, loanNo)`
- `Loan LEFT OUTER JOIN Borrower ON
Loan.loanNo = Borrower.loanNo;`

Questions ??



Thank You !